



# KUBERNETES

( Part 2 )

---

Prepared by: Aamir Pinger



[fb.com/AamirPingerOfficial](https://fb.com/AamirPingerOfficial)



[linkedin.com/in/AamirPinger](https://linkedin.com/in/AamirPinger)



[github.com/AamirPinger](https://github.com/AamirPinger)



# Minikube

- Kubernetes architecture consist of Master node and worker nodes
- Minikube is a tool that makes it easy to run Kubernetes locally
- Minikube runs a single-node Kubernetes cluster on your laptop to use kubernetes for practice or development
- To start Kubernetes cluster using minikube

```
aamir@ap-linux:~$ minikube start
```



```
minikube v1.1.1 on linux (amd64)
```



```
Creating virtualbox VM (CPUs=2, Memory=2048MB, Disk=20000MB) ...
```



```
Configuring environment for Kubernetes v1.14.3 on Docker 18.09.6
```



```
Pulling images ...
```



```
Launching Kubernetes ...
```



```
Verifying: apiserver proxy etcd scheduler controller dns
```



```
Done! kubectl is now configured to use "minikube"
```



# Minikube

- To check the status of minikube

```
aamir@ap-linux:~$ minikube status
```

```
host: Running  
kubelet: Running  
apiserver: Running  
kubectl: Correctly Configured: pointing to minikube-vm at 192.168.99.100
```

- To check addresses of the kubernetes master and services

```
aamir@ap-linux:~$ kubectl cluster-info
```

```
Kubernetes master is running at https://192.168.99.100:8443  
KubeDNS is running at https://192.168.99.100:8443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
```

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

# NODES

---



# Nodes

- Group of server is called a cluster
- Kubernetes cluster is consist of master node and worker nodes
- Master node act as a manager
- Worker nodes are server or system on which kubernetes deploys and run our applications
- To list down nodes in kubernetes cluster

```
aamir@ap-linux:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
Minikube	Ready	master	1m	v1.10.0

You also can use **kubectl get no** command



# Nodes

---

- To check the additional detail of node

```
aamir@ap-linux:~$ kubectl describe nodes minikube
```

- This command show many important things related to worker node, few are
  - Labels
  - Annotations
  - Unschedulable
  - Capacity
  - System Info
  - Allocated resources
  - Events etc



```
aamir@ap-linux:~$ kubectl describe no minikube
Name:          minikube
Roles:         master
Labels:        beta.kubernetes.io/arch=amd64
               beta.kubernetes.io/os=linux
               kubernetes.io/hostname=minikube
               node-role.kubernetes.io/master=
Annotations:   node.alpha.kubernetes.io/ttl: 0
               volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Sat, 01 Dec 2018 17:37:10 +0500
Taints:        <none>
Unschedulable: false
Conditions:
  Type           Status  LastHeartbeatTime   LastTransitionTime   Reason                               Message
  ----           -
  OutOfDisk      False   Wed, 26 Dec 2018 21:44:35 +0500   Sat, 01 Dec 2018 17:37:06 +0500   KubeletHasSufficientDisk           kubelet has sufficient disk space available
  MemoryPressure False   Wed, 26 Dec 2018 21:44:35 +0500   Sat, 01 Dec 2018 17:37:06 +0500   KubeletHasSufficientMemory         kubelet has sufficient memory available
  DiskPressure   False   Wed, 26 Dec 2018 21:44:35 +0500   Sat, 01 Dec 2018 17:37:06 +0500   KubeletHasNoDiskPressure           kubelet has no disk pressure
  PIDPressure    False   Wed, 26 Dec 2018 21:44:35 +0500   Sat, 01 Dec 2018 17:37:06 +0500   KubeletHasSufficientPID            kubelet has sufficient PID available
  Ready          True    Wed, 26 Dec 2018 21:44:35 +0500   Sat, 01 Dec 2018 17:37:06 +0500   KubeletReady                       kubelet is posting ready status

Addresses:
  InternalIP: 10.0.2.15
  Hostname:   minikube
Capacity:
  cpu:                2
  ephemeral-storage:  16888216Ki
  hugepages-2Mi:      0
  memory:             2038624Ki
  pods:              110
Allocatable:
  cpu:                2
  ephemeral-storage:  15564179840
  hugepages-2Mi:      0
  memory:            1936224Ki
  pods:              110
```

ALIAS

---





# Alias

- As you will go further you will see you will be typing kubernetes long command again and again
- We can use Linux alias command to make the long commands shorter. For example

```
aamir@ap-linux:~$ alias kg="kubectl get"
```

```
aamir@ap-linux:~$ kg nodes
```

NAME	STATUS	ROLES	AGE	VERSION
Minikube	Ready	master	1m	v1.14.3

```
aamir@ap-linux:~$ kg no
```

NAME	STATUS	ROLES	AGE	VERSION
Minikube	Ready	master	1m	v1.14.3


PODS

---



# Pods


- To run your application on kubernetes most basic condition is your application should be containerized
- But Kubernetes do not deploy your app's container directly like docker does
- Instead, kubernetes wraps our app container or group of containers together
- This container wrapper is called a Pod





# Pods


- Each pod is like a separate logical machine with its own IP, hostname, processes, and so on, running a single application





# Pods

- All the containers in a pod will appear to be running on the same logical machine
- A Pods with multiple container will always run on a same worker node



WHY PODS ?

---

# ?

Why do we even need pods?

Why can't we use containers directly?

- As we're not supposed to group multiple processes into a single container
- We will be needing another higher-level construct that will allow you to bind containers together and manage them as a single unit

# ?

Why would we even need to run multiple containers together?

- You may have a supporting container with your main application container
- For example a log collection, network proxy etc



# ?

Can't we put all our processes into a single container?

- Yes we can do that, but we don't
- We try to divide our application in smaller microservices
- This makes our app's smaller in size fast to load



# Pods

---

- A pod of containers allows you to run closely related processes together
- Kubernetes provide these containers with (almost) the same environment as if they were all running in a single container, while keeping them somewhat isolated
- “Somewhat isolated”, this is because you want containers inside each group to share certain resources, although not all, so that they’re not fully isolated
- For example if your main application’s container write logs in some file and you want another application to use some part of that log file and make some report



# Pods


---

- Thanks to kubernetes which provides volume concept by which we can share files directories between containers with in the pods
- This is also one of the reason why pods having multiple container does not spread over different nodes. They always co-located on same worker node



# Pods

- Same like dockers, in kubernetes all containers in a Pod also uses same IP but different port numbers
- Port conflicts if same port number allocated only concerns containers in the same pod
- Containers of different pods can never run into port conflicts, because each pod has a separate port space



SHOULD WE RUN  
MULTI-TIER APPS INTO  
SINGLE PODS





# Pods

---

- Yes you can BUT it's not the right way
- For example: If you have a two-node Kubernetes cluster and we have learned multiple container in a single pod always runs on a single worker node
- Disadvantages
  - A single worker node will be utilize
  - Not taking advantage of the computational resources (CPU and memory)
  - You have at your disposal on the second node

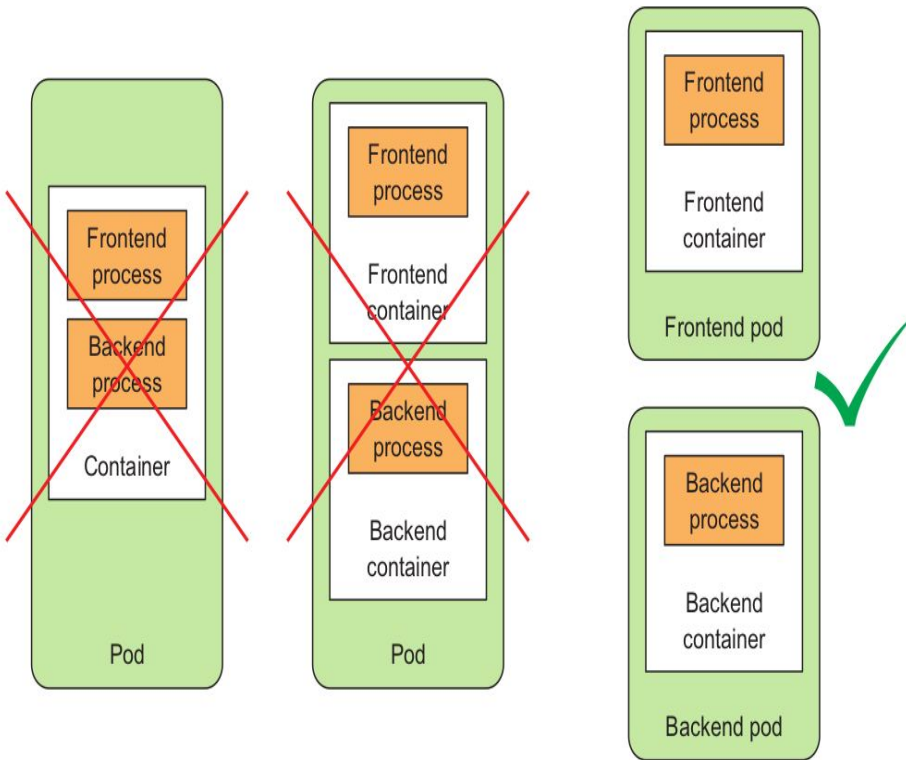
When to put two containers  
on single pod?





# Pods

- When deciding you always need to ask yourself the following questions:
  - Do they need to be run together or can they run on different hosts?
  - Do they represent a single whole or are they independent components?
  - Must they be scaled together or individually?





LET'S CREATE A POD

---



# POD Definition

- To define a POD one way is to define POD configuration in YAML file or JSON format
- **YAML** ("YAML Ain't Markup Language") is commonly used for human-readable configuration files
- **JSON** (JavaScript Object Notation) is a lightweight data-interchange format
- Kuberenet official documentation recommends to write your configuration files using YAML rather than JSON as YAML tends to be more user-friendly

## Sample structure in YAML

```
kind: <Resource Type>
apiVersion: <Kubernete API Version>
metadata:
  name: <Resource name>
spec:
  containers:
    - name: <Container name>
      image: <Container Image>
status:
  containerStatuses:
    ...
  podIP: <Pod IP>
```



# POD Definition

---

- **Kubernetes API version** resource belongs to
  - **Type** of resource
    - Pod
    - Deployments
    - Jobs
    - Others
  - **Metadata** includes the
    - Name
    - Namespace
    - Labels
    - Other information about the pod



# POD Definition

---

- **Spec** contains the actual description of the pod's contents, such as
  - Pod's containers
  - Volumes
  - Security Context
  - Other data
- **Status** contains the current information about the running pod, such as
  - Pod's Condition
  - Every container's description and status
  - Pod's internal IP and other basic info



# Creating Pods with YAML

```
kind: Pod
apiVersion: v1
metadata:
  name: myfirstpod
spec:
  containers:
    - name: container1
      image: aamirpinger/helloworld
      ports:
        - containerPort: 80
```

```
aamir@ap-linux:$ kubectl create -f myfirstpod.yaml
Pod "myfirstpod" created
```

# POD Listing

---



# Pod Listing

---

```
aamir@ap-linux:$ kubectl get pods
```

```
Pod "myfirstpod" created
```

NAME	READY	STATUS	RESTARTS	AGE
myfirstpod	1/1	Running	0	15m

You can even write **kubectl get po**

# YAML FROM POD

---





# YAML FROM POD

- We created pod with YAML, What if we have already a pod and we wanted to check pod insights

```
aamir@ap-linux:$ kubectl get pod myfirstpod -o yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2019-07-04T06:36:20Z"
  name: myfirstpod
  namespace: default
  resourceVersion: "1994"
  selfLink: /api/v1/namespaces/default/pods/myfirstpod
  uid: 08b74341-9e26-11e9-b5e1-080027713168
spec:
  containers:
  - image: aamirpinger/helloworld
  ...
```

# POD CREATION WITHOUT YAML

---



# Pod Creation Without Yaml

- To create a pod without writing YAML file we can use

**Kubectrl run <pod\_name> --image=<image\_name>  
--port=<container\_exposed\_port> --restart=Never**

```
aamir@ap-linux:$ kubectl run mysecondpod  
--image=aamirpinger/helloworld --port=80 --restart=Never  
pod/mysecondpod created
```

```
aamir@ap-linux:$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
myfirstpod	1/1	Running	0	70m
mysecondpod	0/1	ContainerCreating	0	4s

# PORT FORWARDING

---



# Port Forwarding

---

- Our pod is created which got our helloworld application but still not accessible at browser

**kubectl port-forward** <pod\_name> <external\_port>:<container\_port>

```
aamir@ap-linux:$ kubectl port-forward mysecondpod 6500:80
Forwarding from 127.0.0.1:6500 -> 80
Forwarding from [::1]:6500 -> 80
```

<http://localhost:6500>

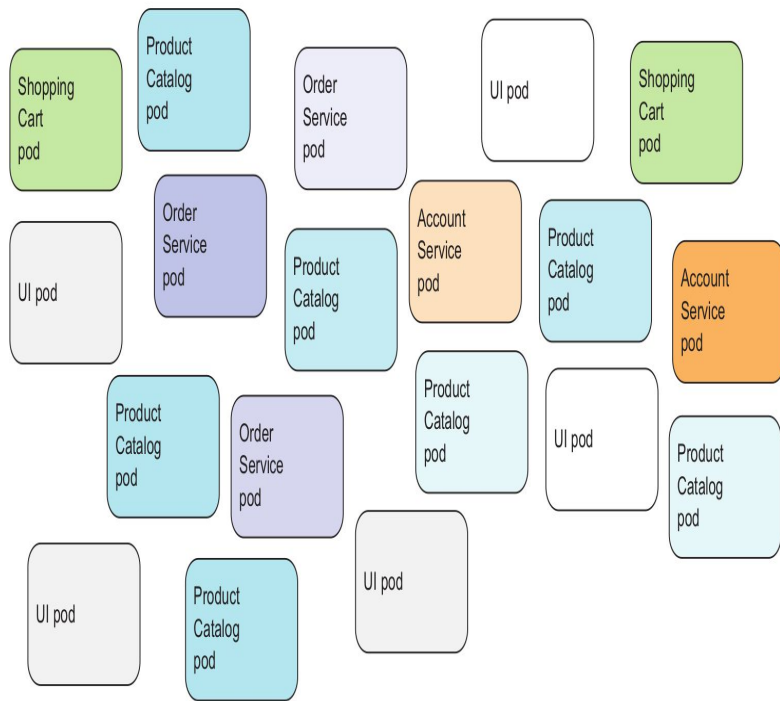
# LABELS

---



# Labels

- Uptil now we have successfully create a pod having one container
- Theoretically, we have learnt that kubernetes does help creating  $n^{\text{th}}$  number of replicas
- Theoretically, we have also learnt that kubernetes can groups pods providing at a single static IP address
- Think of situation where you have hundreds of pods with multiple replicas





Question  
arises

What mechanism  
kubernetes use for  
organizing them?





Question  
arises

How group of pods can  
be managed with a  
single action instead of  
having to perform the  
action for each pod  
individually



# Labels

---

- Kubernetes organizes pods with the help of labels
- Labels are simple key/value pair that kubernetes uses with almost all the resources it creates to group them
- Kubernetes gives option to filter the resources by providing label criteria
- Multiple labels can be assigned to any resource and they become part of resource group having with same labels
- Labels can be assigned at a time of resource creation or even you can assign or change label value after creation of the resource



# Creating Pods With Labels

```
kind: Pod
apiVersion: v1
metadata:
  name: myfirstpodwithlabel
```

```
  labels:
    type: backend
    env: production
```

```
spec:
  containers:
```

- name: container1
- image: aamirpinger/helloworld
- ports:
- containerPort: 80

```
aamir@ap-linux:$ kubectl create -f myfirstpodwithlabel.yaml
Pod "myfirstpodwithlabel" created
```



# Creating Pods With Labels

- To create a pod without writing YAML file we can use

```
Kubectl run <pod_name> --image=<image_name> --port=<container_exposed_port>  
--restart=Never --labels=key1=val1,key2=val2...,keyN=valN
```

```
aamir@ap-linux:$ kubectl run anotherpodwithlabel  
--image=aamirpinger/helloworld --port 80 --restart=Never  
--labels=type=frontend,env=develop  
pod/anotherpodwithlabel created
```

# PODS LISTING WITH LABELS

---



# Pods Listing With Labels

- To get the labels in the listing

```
aamir@ap-linux:$ kubectl get po --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
myfirstpod	1/1	Running	0	70m	<none>
anotherpodwithlabel	0/1	Running	0	4s	env=develop,type=frontend
...					

- To make label as a column

```
aamir@ap-linux:$ kubectl get po -L env,type
```

NAME	READY	STATUS	RESTARTS	AGE	ENV	TYPE
myfirstpod	1/1	Running	0	70m		
myfirstpodwithlabel	1/1	Running	0	7m8s	production	
anotherpodwithlabel	0/1	Running	0	4s	develop	frontend
...						

# PODS LABELS AT RUNTIME

---



# Pods Listing With Labels

- Assigning labels to existing pods

```
aamir@ap-linux:$ kubectl label pod myfirstpod app=helloworld type=frontend
```

```
aamir@ap-linux:$ kubectl get po --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
myfirstpod	1/1	Running	0	70m	app=helloworld,type=frontend
anotherpodwithlabel	1/1	Running	0	4s	env=develop,type=frontend

- Modifying existing labels of existing pods

```
aamir@ap-linux:$ kubectl label pod anotherpodwithlabel env=prod --overwrite
```

```
aamir@ap-linux:$ kubectl get po --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
myfirstpod	1/1	Running	0	71m	app=helloworld,type=frontend
anotherpodwithlabel	1/1	Running	0	1m	env=prod,type=frontend

...



# PODS LISTING WITH LABEL SELECTOR

---



# Pods Listing With Label Selector

---

- Label selector can be used as criteria for filtering any resources
- We can use label as criteria to check if
  - Contains (or doesn't contain) a label with a certain key
  - Contains a label with a certain key and value
  - Contains a label with a certain key, but with a value not equal to the one you specify
  - Contains (or doesn't contain) any one of supplied value in any particular label key



# Pods Listing With Labels

- Contains a label with a certain key/value

```
aamir@ap-linux:$ kubectl get po -l type=frontend
```

NAME	READY	STATUS	RESTARTS	AGE
myfirstpod	1/1	Running	0	3h59m
anotherpodwithlabel	1/1	Running	0	36m

...

```
aamir@ap-linux:$ kubectl get po -l type=frontend --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
myfirstpod	1/1	Running	0	4h1m	app=helloworld,type=frontend
anotherpodwithlabel	1/1	Running	0	38m	env=develop,type=frontend

...



# Pods Listing With Labels

- Contains a multiple labels with a certain key/value

```
aamir@ap-linux:$ kubectl get po -l env=prod,type=frontend --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
anotherpodwithlabel	1/1	Running	0	53m	env=prod,type=frontend

- Does NOT contains a label value against certain key

```
aamir@ap-linux:$ kubectl get po -l type!=frontend --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
myfirstpodwithlabel	1/1	Running	0	37m	type=backend,env=production
mysecondpod	1/1	Running	0	138m	run=mysecondpod

...



# Pods Listing With Labels

- Does contains a label irrespective of the value

```
aamir@ap-linux:$ kubectl get po -l type --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
myfirstpod	1/1	Running	0	4h10m	app=helloworld,type=frontend
anotherpodwithlabel	1/1	Running	0	47m	env=prod,type=frontend

```
...
```

- Does NOT contains a label irrespective of the value

```
aamir@ap-linux:$ kubectl get po -l '!type' --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
mysecondpod	1/1	Running	0	138m	run=mysecondpod



# Pods Listing With Labels

- Contains any value from provided list of values against any key

```
aamir@ap-linux:$ kubectl get po -l 'type in (backend,frontend)' --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
myfirstpod	1/1	Running	0	4h29m	app=helloworld,type=frontend
myfirstpodwithlabel	1/1	Running	0	62m	env=production,type=backend
anotherpodwithlabel	1/1	Running	0	53m	env=prod,type=frontend

- Does NOT contains any value from provided list of values against any key

```
aamir@ap-linux:$ kubectl get po -l 'type notin(backend,frontend)' --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
mysecondpod	1/1	Running	0	165m	run=mysecondpod

# PODS SCHEDULING WITH NODE SELECTOR

---



# Pods Scheduling With Node Selector

```
kind: Pod
apiVersion: v1
metadata:
  name: podwithnodeselector
spec:
  nodeSelector:
    disk_type: "ssd"
  containers:
    - name: container1
      image: aamirpinger/helloworld
      ports:
        - containerPort: 80
```

This will deploy this new pod on a worker node where label is **disk\_type=ssd**



# ANNOTATION

---



# Annotation

---

- Kubernetes also provide annotation feature with labels
- Annotation are basically words that explanation or comment on something
- For example, annotation can be used to write creator's name or contact or about application it running
- Labels are meant to hold limited information whereas you can have larger annotation with any resource
- Annotation are way to add some extra information to the resource but it cannot be used to group or filter like labels



# Pod Creation With Annotation

```
kind: Pod
apiVersion: v1
metadata:
  name: podwithannotation
  annotations:
    createdBy: Aamir Pinger
    email: aamirpinger@yahoo.com
spec:
  containers:
    - name: container1
      image: aamirpinger/helloworld
      ports:
        - containerPort: 80
```

```
aamir@ap-linux:$ kubectl create -f
podwithannotation.yaml
Pod "podwithannotation" created
```



# Annotation Assignment To Any Resource

- To annotate any existing resource

```
aamir@ap-linux:$ kubectl annotate pod myfirstpod createdBy="Aamir  
Pinger" email="aamirpinger@yahoo.com"  
pod/myfirstpod annotated
```

- Let's check by re generating YAML from pod

```
aamir@ap-linux:$ kubectl get pod myfirstpod -o yaml  
email="aamirpinger@yahoo.com"  
apiVersion: v1  
kind: Pod  
metadata:  
  annotations:  
    createdBy: Aamir Pinger  
    email: aamirpinger@yahoo.com
```

# DESCRIBING POD'S INSIGHTS

---



# DESCRIBING POD'S INSIGHTS

- We can have any resource insight with `kubectl describe` command

```
aamir@ap-linux:$ kubectl describe pod myfirstpod
```

```
Name:          myfirstpod
Namespace:     default
Priority:       0
PriorityClassName: <none>
Node:          minikube/10.0.2.15
Start Time:    Thu, 04 Jul 2019 11:36:20 +0500
Labels:        app=helloworld
               type=frontend
Annotations:   createdBy: Aamir Pinger
               email: aamirpinger@yahoo.com
Status:        Running
IP:            172.17.0.7
Containers:
  container1:
```

```
...
```

# NAMESPACE

---



# Namespace

---

- With labels we have seen how easy is to group resources
- But what if any label overlaps, for example
  - `type=frontend` is assigned to few pods with `app_name=web` and `env=development`
  - `type=frontend` is assigned to few pods with `app_name=web` and `env=production`
- Name of particular resource will always be unique, what about setting the similar environment as production when in development, QA phase
- But what about times when you want to split objects into separate, non-overlapping groups? You may want to only operate inside one group at a time





# Namespace

---

- For this and other reasons, Kubernetes also groups objects into namespaces
- Namespace is a kind of virtualbox which isolate self contain resources with other namespace
- You can easily separate the scope of resources using namespaces. E.g. Resource in namespace for development phase cannot harm resources in namespace for production
- Similarly you can split them into multiple namespaces, which also allows you to use the same resource names multiple times (across different namespaces)



# Namespace

- To create a namespace

```
aamir@ap-linux:$ kubectl create namespace production
namespace/production created
```

**Kubectl create ns production** can also be used

- To get a list of namespaces

```
aamir@ap-linux:$ kubectl get namespace
```

NAME	STATUS	AGE
default	Active	37h
kube-node-lease	Active	37h
kube-public	Active	37h
kube-system	Active	37h
production	Active	51s



# Namespace

- To create a pod in any particular namespace

```
aamir@ap-linux:$ kubectl run myfirstpod --image aamirpinger/helloworld  
--restart=Never --namespace=production  
pod/myfirstpod created
```

- To get any resource from any particular namespace

```
aamir@ap-linux:$ kubectl get pod --namespace production
```

NAME	READY	STATUS	RESTARTS	AGE
myfirstpod	1/1	Running	0	4m45s

- To get any resource from all namespaces

```
aamir@ap-linux:$ kubectl get pod --all-namespace
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	myfirstpod	1/1	Running	0	36h ...

# DELETING A RESOURCE

---



# DELETING A RESOURCE

**Kubectrl delete** <resource\_type> <resource\_name>  
**--namespace=**<optional\_namespace\_name>

- To delete pod from default namespace

```
aamir@ap-linux:$ kubectl delete pod myfirstpod  
Pod "myfirstpod" deleted
```

- To delete pod from from any particular namespace

```
aamir@ap-linux:$ kubectl delete pod myfirstpod --namespace production  
Pod "myfirstpod" deleted
```



# DELETING A RESOURCE

- To delete pod based on label criteria

```
aamir@ap-linux:$ kubectl delete pod -l kubectl delete pod -l 'type  
in(frontend,backend)'  
pod "anotherpodwithlabel" deleted  
pod "myfirstpodwithlabel" deleted
```

- To delete all pods

```
aamir@ap-linux:$ kubectl delete pod --all  
pod "anotherpodwithlabel" deleted  
pod "myfirstpodwithlabel" deleted
```

- To delete pod from from any particular namespace

```
aamir@ap-linux:$ kubectl delete ns production  
namespace "production" deleted
```

# ReplicaSets

---



# ReplicaSets

---

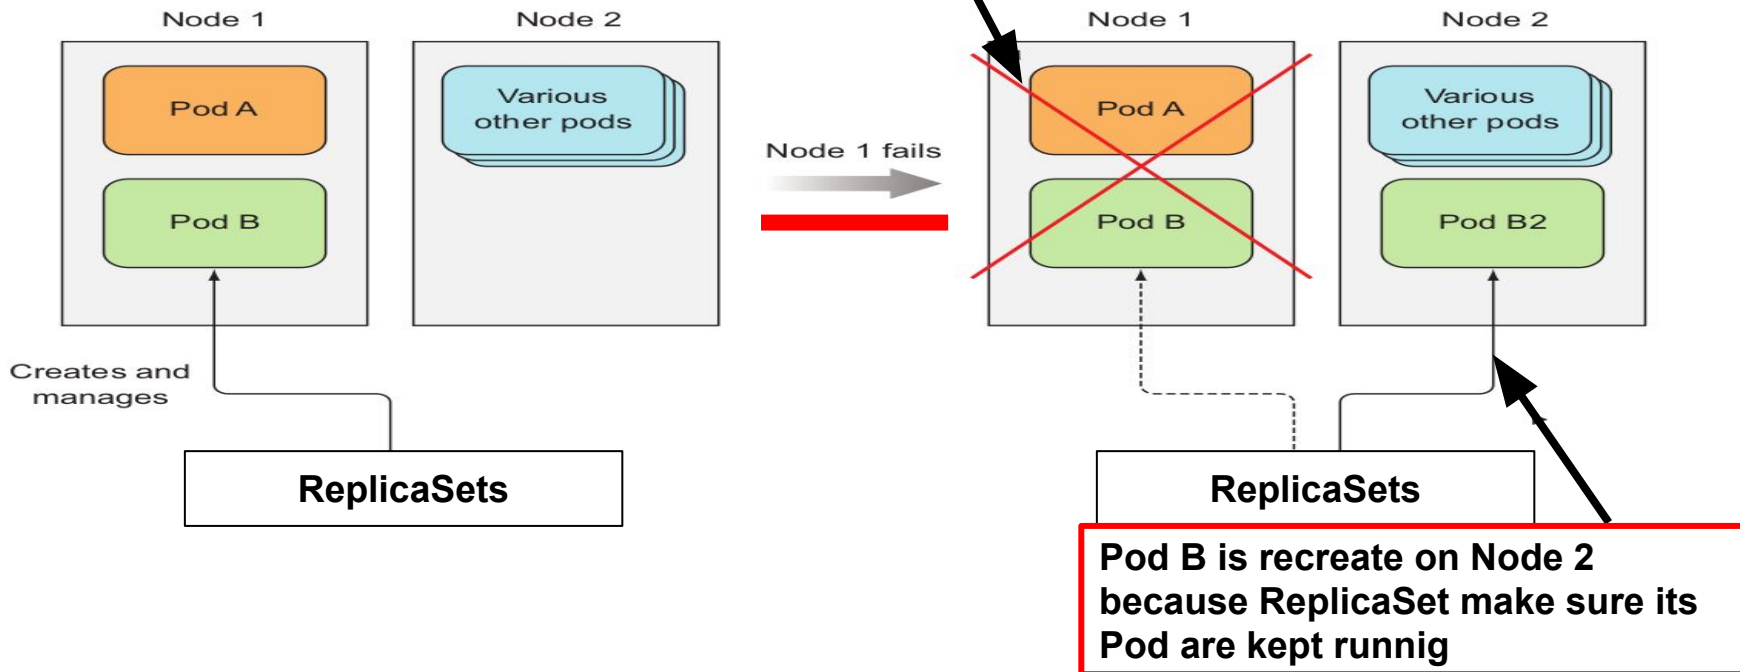
- ReplicaSet is also one of the resource in kubernetes like pods and others are
- ReplicaSet is the resource that help creating and managing multiple copies of application (replicas) in kubernetes
- When we create any pod through ReplicaSet It ensures its pods are always kept running
- If the pod disappears for any reason, such as
  - If the worker node disappears from the cluster
  - If the pod was evicted from the worker node
- The ReplicaSet notices the missing pod and creates a replacement pod





# ReplicaSets

**Pod A goes down with Node 1 and not recreated, because it was not created with ReplicaSets**





# ReplicaSets

---

- A ReplicaSet has three essential parts
  1. Label selector
    - Determines what pods are in the ReplicaSet scope
  2. Replica count
    - Specifies the desired number of pods that should be running
  3. Pod template
    - ReplicaSet uses to create new pod



# ReplicaSets

## replicaset.yaml

```
apiVersion: apps/v1beta2
```

```
kind: ReplicaSet
```

```
metadata:
```

```
name: my-replica-set
```

```
spec:
```

```
replicas: 3
```

```
selector:
```

```
  matchLabels:
```

```
    app: myapp
```

```
template:
```

```
  metadata:
```

```
    labels:
```

```
      app: myapp
```

These two must  
match, otherwise  
infinite number  
of pod will be  
created

```
metadata:
```

```
  ...
```

```
spec:
```

```
  containers:
```

```
    - name: my-rs-container
```

```
      image: aamirpinger/flag:latest
```

```
      ports:
```

```
        - containerPort: 80
```

```
aamir@ap-linux:~$ kubectl create -f  
replicaset.yaml  
replicaset.apps/my-replica-set created
```



# ReplicaSets

- To get the list of ReplicaSets resource

```
aamir@ap-linux:~$ kubectl get replicaset
```

NAME	DESIRED	CURRENT	READY	AGE
my-replica-set	3	3	3	5m6s

- **Kubectl get rs can also be used**
- To get the list of ReplicaSets resource

```
aamir@ap-linux:~$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
my-replica-set-m644t	1/1	Running	0	8m55s
my-replica-set-4m4kt	1/1	Running	0	8m55s
my-replica-set-njndl	1/1	Running	0	8m55s
....				



# ReplicaSets

- Let's manually delete one pod and check what effect it gets on pods listing

```
aamir@ap-linux:~$ kubectl delete pod my-replica-set-m644t
pod "my-replica-set-m644t" deleted
```

- Let's get the list of ReplicaSet and pods together

```
aamir@ap-linux:~$ kubectl get rs,pod
```

NAME	READY	STATUS	RESTARTS	AGE
pod/my-replica-set-4m4kt	1/1	Running	0	8m55s
pod/my-replica-set-njndl	1/1	Running	0	8m55s
pod/my-replica-set-snhjt	0/1	ContainerCreating	0	3s

....

NAME	DESIRED	CURRENT	READY	AGE
replicaset.extensions/my-replica-set	3	3	2	13m



# ReplicaSets

- Let's manually delete one pod and check what effect it gets on pods listing

```
aamir@ap-linux:~$ kubectl describe rs my-replica-set
```

```
Name:      my-replica-set
Namespace:  default
Selector:   app=myapp
Labels:     <none>
Annotations: <none>
Replicas:   3 current / 3 desired
Pods Status: 3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=myapp
  Containers:
    my-rs-container:
      Image:      aamirpinger/flag:latest
      ...
```



# ReplicaSets

- Let's make a Pod out of scope of ReplicaSets by changing label of any pod

```
aamir@ap-linux:~$ kubectl label pod my-replica-set-njndl app=myappNEW --overwrite
pod/my-replica-set-njndl labeled
```

- Let's get the list of ReplicaSet and pods together

```
aamir@ap-linux:~$ kubectl get rs,pod --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
pod/my-replica-set-4m4kt	1/1	Running	0	18m	app=myapp
pod/my-replica-set-8xtwv	0/1	<u>ContainerCreating</u>	0	4s	app=myapp
pod/my-replica-set-njndl	1/1	Running	0	18m	app=myappNEW
pod/my-replica-set-snhjt	1/1	Running	0	13m	app=myapp

....

NAME	DESIRED	CURRENT	READY	AGE	LABELS
replicaset.extensions/my-replica-set	3	3	2	13m	<none>



# ReplicaSets

- To edit any resource you can also use  
**kubectl edit <resource type> <resource name>**
- Let's make the label again app=myapp of pod my-replica-set-njndl

```
aamir@ap-linux:~$ kubectl edit pod my-replica-set-njndl
```

```
pod/my-replica-set-njndl edited
```

```
aamir@ap-linux:~$ kubectl get rs,pod --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
pod/my-replica-set-4m4kt	1/1	Running	0	26m	app=myapp
pod/my-replica-set-8xtwv	0/1	<u>Terminating</u>	0	10m	app=myapp
pod/my-replica-set-njndl	1/1	Running	0	26m	app=myapp
pod/my-replica-set-snhjt	1/1	Running	0	20m	app=myapp

```
....
```

NAME	DESIRED	CURRENT	READY	AGE	LABELS
replicaset.extensions/my-replica-set	3	3	3	26m	<none>





# ReplicaSets

- We used matchLabels in selector of ReplicaSet configuration in which we can add multiple labels
- Based on those labels ReplicaSet will look for pods that have all those labels together and group them
- In ReplicaSet you can even add additional expressions to the selector

```
...
metadata:
name: my-replica-set
spec:
replicas: 3
selector:
  matchExpressions:
    - key: app
      operator: In
      values:
        - myapp
template:
  metadata:
    labels:
...

```



# ReplicaSets

- There are 4 operators for matchExpressions
  - **In,**
  - **notIn,**
  - **Exists,**
  - **DoesNotExist**
- It's important to remember that If you specify both matchLabels and matchExpressions then all the labels must match to group the resource



# ReplicaSets

---

- All of the three parts of ReplicaSet can be modified at runtime
- Only changes to the replica count will affect existing pods
- Changes to label selector and pod template in ReplicaSet will only affect new containers
- Due to change in label selector or pod label itself, If existing pods fall out of the scope of the ReplicaSet, so the controller stops caring about them

# POD SCALING

---



# Pod Scaling

- ReplicaSets make sure a desired number of pod instances is always running
- Scaling number of pods up and down can be done anytime
- If you will scale up, ReplicaSet will add more pod
- If you will scale down, ReplicaSet will terminate pods to match the numbers

```
aamir@ap-linux:~$ kubectl scale rs my-replica-set --replicas=5
```

```
replicaset.extensions/my-replica-set scaled
```

```
aamir@ap-linux:~$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
replicaset.extensions/my-replica-set	5	5	5	4h17m

```
aamir@ap-linux:~$ kubectl scale rs my-replica-set --replicas=2
```

```
replicaset.extensions/my-replica-set scaled
```



# Pod Scaling

- Even if you delete or add a pod manually with the same labels used by replicaSets to group the pods, it will automatically add or terminate pods to match the exact numbers provided to replicaSets
- When you delete a replicaSet all the pods under that replicaSet will also get terminated
- Kubernetes does provide `--cascade=false` option for deleting only replicaSet and not the pods under it

```
aamir@ap-linux:~$ kubectl delete rs my-replica-set --cascade=false  
replicaset.extensions "my-replica-set" deleted
```

JOB

---



# Job

---

- Job is another resource in kubernetes
- It's basically a Pod which we create under the type (kind) of Job
- A job resource is used to create a pod which terminates automatically when the defined job of that pods successfully completed
- A job resource creates one or more pods and ensures that a specified number of them successfully terminate





# Job

---

- The Job object will start a new Pod if the first pod fails or is deleted (for example due to a node hardware failure or a node reboot)
- A Job can also be used to run multiple pods in parallel
- Deleting a Job resource will cleanup the pods it created



# Job

```
apiVersion: batch/v1
kind: Job
metadata:
  name: whalesay
spec:
  template:
    spec:
      containers:
      - name: whalesay
        image: docker/whalesay
        command: ["cowsay", "This is a Kubernetes Job!"]
        restartPolicy: Never
  backoffLimit: 4
  activeDeadlineSeconds: 60
```

**restartPolicy:** **Never** or **OnFailure** is allowed

- **OnFailure** - the Pod stays on the node, but the Container get a restart
- **Never** - the Job controller starts a new Pod and leave the unsuccessful pod as it is



# Job

```
apiVersion: batch/v1
kind: Job
metadata:
  name: whalesay
spec:
  template:
    spec:
      containers:
      - name: whalesay
        image: docker/whalesay
        command: ["cowsay", "This is a Kubernetes Job!"]
        restartPolicy: Never
  backoffLimit: 4
  activeDeadlineSeconds: 60
```

## **backoffLimit**

- Specify the number of retries before considering a Job as failed
- 6 is default retries set back-off limit by kubernetes
- Failed Pods associated with the Job are recreated by the Job controller with an exponential back-off delay (10s, 20s, 40s ...) capped at six minutes



# Job

```
apiVersion: batch/v1
kind: Job
metadata:
  name: whalesay
spec:
  template:
    spec:
      containers:
      - name: whalesay
        image: docker/whalesay
        command: ["cowsay", "This is a Kubernetes Job!"]
        restartPolicy: Never
      backoffLimit: 4
      activeDeadlineSeconds: 60
```

## **activeDeadlineSeconds**

- Applies to the duration of the job
- Once a Job reaches activeDeadlineSeconds, all of its Pods are terminated and the Job status will become type: Failed with reason: DeadlineExceeded



# Job

```
aamir@ap-linux:~$ kubectl create -f job.yaml
```

```
job.batch/whalesay created
```

```
aamir@ap-linux:~$ kubectl get po,job
```

NAME	READY	STATUS	RESTARTS	AGE
pod/whalesay-qgsh7	0/1	Completed	0	10s

```
...
```

NAME	COMPLETIONS	DURATION	AGE
job.batch/whalesay	1/1	6s	10s

```
aamir@ap-linux:~$ kubectl describe job whalesay
```

```
Name:          whalesay
```

```
...
```

```
Active Deadline Seconds: 60s
```

```
Pods Statuses:      0 Running / 1 Succeeded / 0 Failed
```

```
Pod Template:
```

```
...
```



< This is a Kubernetes Job! >

Source: Kubernetes in Action Book by Marko Luksa (Manning Publications)

# CRONJOB

---



# CronJob

---

- This works almost similar to Kubernetes Job resource
- Only difference is Job resource create Pod instantly and once job is completed it does not recreate Job resource or pod
- CronJob resource is used to schedule Job at later time and can be set to initiate the Job again on provided time gap
- CronJob always creates only a single Job resource for each execution configured in the schedule





# CronJob

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: batch-job-every-minute
spec:
  schedule: "* * * * *"
  jobTemplate:
    spec:
      template:
        metadata:
          labels:
            app: periodic-batch-job
        spec:
          restartPolicy: OnFailure
          containers:
            - name: main
              image: docker/whalesay
              command: ["cowsay", "This is a CronJob!"]
```

```
aamir@ap-linux:~$ kubectl create -f cronjob.yaml
cronjob.batch/batch-job-every-minute created
```



# CronJob

```
aamir@ap-linux:~$ kubectl create -f cronjob.yaml
```

```
cronjob.batch/batch-job-every-minute created
```

```
aamir@ap-linux:~$ kubectl get po,cj
```

NAME	READY	STATUS	RESTARTS	AGE
pod/batch-job-every-minute-1562696940-fznlm	0/1	Completed	0	30s

```
...
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
cronjob.batch/batch-job-every-minute	* * * * *	False	0	38s	50s

```
aamir@ap-linux:~$ kubectl describe cj batch-job-every-minute
```

```
Name:          batch-job-every-minute
Namespace:     default
Labels:        <none>
Annotations:   <none>
Schedule:      * * * * *
Concurrency Policy:  Allow
```

```
...
```




< This is a Kubernetes CronJob! >

```
aamir@ap-linux:~$ kubectl delete cj batch-job-every-minute
```



# CronJob

## Schedule pattern



## Examples:

Every Minute

\* \* \* \* \*

Everyday 5am and 5pm

0 5,17 \* \* \*

Every midnight in weekdays

0 0 \* \* 1-5

Every 15 minutes

\*/15 \* \* \* \*

Every hour of alternate days

0 \* \*/2 \* \*

# Thank you and God bless you all!

---



[fb.com/AamirPingerOfficial](https://fb.com/AamirPingerOfficial)



[linkedin.com/in/AamirPinger](https://linkedin.com/in/AamirPinger)



[github.com/AamirPinger](https://github.com/AamirPinger)